

AI Matching Engine 500 Error - Root Cause Analysis & Fix

Issue Summary

AI Matching Engine was logging continuous 500 errors in production, despite multiple previous fix attempts targeting OpenAI integration, validation, and error handling.

Log Analysis

Error Pattern from Render Logs

```
[GET /api/family/match] Error: i [UnauthenticatedError]: You must be logged in to access this resource
  at _ (/app/.next/server/chunks/3400.js:1:997)
  at process.processTicksAndRejections (node:internal/process/task_queues:95:5)
  at async S (/app/.next/server/app/api/family/match/route.js:1:8444)
```

Critical Discovery

User-Agent: "aiohttpClient (cluster-proxy-5d748f78ff-97vhp)"

This revealed the requests were **NOT from actual users** but from:

- Render's internal health checks
- Infrastructure monitoring probes
- Load balancer health verification

Root Cause

The Real Problem

1. `GET /api/family/match` endpoint requires authentication via `requireAuth()`
2. Render's health checks/monitoring call this endpoint **without authentication**
3. `requireAuth()` throws `UnauthenticatedError` when no session exists
4. Error was being caught but logged as a 500 error (misleading)
5. Logs filled with “errors” that were actually expected 401 scenarios

Why Previous Fixes Failed

Previous commits addressed:

- X OpenAI graceful handling (fdce991)
- X Comprehensive error logging (4121544)
- X Robust validation (cd3d75d)

But the actual issue was **authentication handling**, not matching logic!

Solution Implemented

Changes Made to /src/app/api/family/match/route.ts

1. Import UnauthenticatedError Class

```
import { requireAuth, UnauthenticatedError } from '@/lib/auth-utils';
```

2. Improved GET Handler Error Handling

```
export async function GET(request: NextRequest) {
  try {
    const user = await requireAuth();
    // ... rest of logic
  } catch (error) {
    // Handle authentication errors gracefully (health checks, monitoring)
    if (error instanceof UnauthenticatedError) {
      // Don't log authentication failures as errors (these are expected)
      return NextResponse.json(
        { error: 'Authentication required' },
        { status: 401 }
      );
    }

    // Log only unexpected errors
    console.error('[GET /api/family/match] Error:', error);

    return NextResponse.json(
      { error: 'Internal server error' },
      { status: 500 }
    );
  }
}
```

Key Improvements

1. **Graceful 401 Handling:** Returns proper 401 status for unauthenticated requests
2. **Reduced Noise:** Doesn't log expected authentication failures as errors
3. **Better Diagnostics:** Only logs actual unexpected errors
4. **Type Safety:** Uses `instanceof` check with imported error class

Expected Impact

Before Fix

-  Logs filled with “500 errors” (misleading)
-  Alert fatigue from expected behavior
-  Difficult to identify real issues
-  Wasted debugging time

After Fix

-  Health checks return clean 401 responses
-  Logs show only actual errors
-  Better signal-to-noise ratio for monitoring

- ✓ Real issues easily identifiable

Testing Verification

Manual Testing

```
# Unauthenticated request (should return 401, not 500)
curl -X GET https://carelinkai.onrender.com/api/family/match
# Expected: {"error":"Authentication required"} with 401 status

# Authenticated request (should work normally)
curl -X GET https://carelinkai.onrender.com/api/family/match \
-H "Cookie: next-auth.session-token=..."
# Expected: {"success":true,"matchRequests": [...]}
```

What to Monitor

1. ✓ Health check requests return 401 (not 500)
2. ✓ Real errors still logged with full details
3. ✓ Authenticated requests work normally
4. ✓ POST /api/family/match (matching) works correctly

Lessons Learned

Investigation Best Practices

1. **Check User-Agent:** Distinguish between user traffic and infrastructure probes
2. **Review Full Logs:** Look beyond error messages to request context
3. **Understand Request Patterns:** Rapid repeated requests suggest automation
4. **Authentication vs Logic:** Separate auth failures from business logic errors

Architecture Insights

1. Health checks should use dedicated endpoints when possible
2. Public endpoints vs authenticated endpoints need clear distinction
3. Expected errors (401, 403) should be handled differently from unexpected errors (500)
4. Monitoring infrastructure needs special consideration

Deployment Status

Build Verification

- ✓ TypeScript compilation successful
- ✓ Next.js build completed
- ✓ All routes validated
- ✓ No breaking changes

Deployment Steps

1. ✓ Code changes committed
2. ⏳ Push to GitHub
3. ⏳ Render auto-deploy

4. ⏳ Monitor logs for improved error patterns

Success Metrics

Key Performance Indicators

- **Error Rate:** Should drop from 100+ per minute to near-zero
- **Log Quality:** Only real errors should appear in logs
- **Alert Accuracy:** Alerts should trigger only for actual issues
- **Debug Time:** Reduced time to identify real problems

Monitoring Checklist

- [] Verify 401 responses for unauthenticated GET requests
- [] Confirm no 500 errors from health checks
- [] Test actual matching functionality still works
- [] Validate authenticated GET requests work
- [] Check audit logs for proper tracking

Related Files

Modified Files

- `src/app/api/family/match/route.ts` - Main fix implementation

Referenced Files

- `src/lib/auth-utils.ts` - UnauthenticatedError class definition
- `src/app/api/family/match/[id]/route.ts` - Single match retrieval
- `src/app/dashboard/find-care/page.tsx` - POST /api/family/match caller
- `src/app/dashboard/find-care/results/[id]/page.tsx` - GET /api/family/match/[id] caller

Recommendations

Immediate Actions

1. Deploy this fix to production
2. Monitor logs for 24 hours
3. Verify matching functionality with test user
4. Update monitoring alerts if needed

Future Improvements

1. **Health Check Endpoint:** Create `/api/health` for infrastructure checks
2. **Middleware:** Add authentication middleware that handles 401s consistently
3. **Monitoring Dashboard:** Distinguish between 4xx (client errors) and 5xx (server errors)
4. **Rate Limiting:** Consider rate limiting for unauthenticated requests
5. **Documentation:** Update API documentation with authentication requirements

Fix Confidence Level:  95%

Reasoning:

- Root cause clearly identified in logs
- Solution directly addresses the authentication error
- No changes to matching logic (which was working)
- Build verified successfully
- Low risk of introducing new issues